



WHAT EVEN IS REINFORCEMENT LEARNING?




Reinforcement Learning enables intelligent agents to take actions to maximize rewards




You don't need labels and you don't need corrections to suboptimal actions



You balance exploration with exploitation. Try processes!



Using time steps, choose an action to get a reward. Transition, choose a new state, get a new reward



The goal is to learn a policy to maximize cumulative reward

Common algorithms: multi armed bandit, monte carlo methods

